



assignment2

Updated automatically every 5 minutes

## Assignment #2 CS695 Spring 2023-24

### Topic: Enter the VM

#### Statutory note:

- This course is on a ***no-plagiarism*** diet.
- All parties involved in plagiarism harakiri will be penalized to the maximum extent.
- The Moss Detective Agency has agreed to conduct all investigations.  
<https://theory.stanford.edu/~aiken/moss/>
- Byomkesh, Sherlock, Phryne, Marple, and Hercule are on standby.
- Hardcoding the runtime output in the code will be heavily penalized.
- Generative AI (ChatGPT, Gemini, etc.) is your friend, but it cannot generate outputs for you to submit. If you want to submit such generated outputs, mark them explicitly in red in your submissions.
- **Warning: Submission Guidelines should be strictly followed; otherwise, your submission will not count. (0 Marks)**

## 0. Introduction

Virtualization techniques are



## assignment2

Updated automatically every 5 minutes

of VMM design (Popek & Goldberg [1]). This is because x86 is very difficult to virtualize. Thus, hardware vendors now support x86 hardware virtualization natively.

KVM or Kernel-based Virtual Machine allows you to create hypervisors in Linux, which can be controlled by userspace programs to run guest VMs while handling x86 hardware virtualization internally. KVM is implemented as a Linux kernel module and exposes a device (`/dev/kvm`) that hypervisors can use to perform operations through `ioctl()` calls.

KVM API allows userspace hypervisors to perform the following operations:

- Creation of new virtual machines
- Allocation of memory to virtual machines.
- Reading and writing virtual CPU registers.
- Injecting and interrupting into a virtual CPU.
- Running a virtual CPU.

References:

- KVM Paper: [\[link1\]](#)
- Linux KVM API: [\[link2\]](#) [\[link3\]](#)

The userspace program relies on the KVM hypervisor for the x86 architecture virtualization. Still, it has to implement its own IO handling and device support — block IO, network IO, console drivers, USB controller, host file system, etc. In practice, QEMU is used as a userspace program alongside KVM, which implements (emulates) all these components and can run any OS as a guest.

**Setup Procedure (for Ubuntu):**  
<https://help.ubuntu.com/>



Published using Google Docs

[Report abuse](#)[Learn more](#)

assignment2

Updated automatically every 5 minutes

**Download the Assignment code tarball:** [link](#)

**Bonus:** You can try out QEMU + KVM to run any x86-64 OS of your choice in Linux KVM <https://ubuntu.com/server/docs/virtualization-virt-tools>

## 1a. DIY Hypervisor

In this assignment, you will build your VM using KVM from scratch.

*a lot of boilerplate code with a simple hypervisor implemented.*

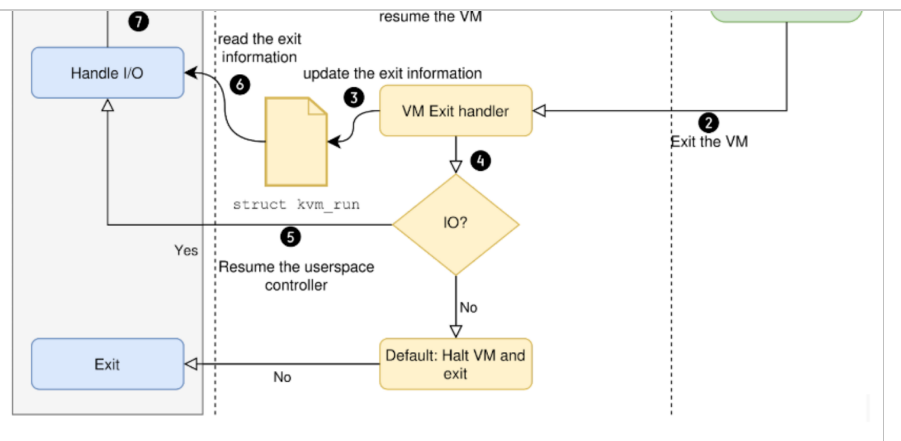
Starter Code: `assignment2/part1/`

The simple-kvm hypervisor is implemented in the source file `simple-kvm.c`. Since your hypervisor is simple, you cannot run a full-fledged OS. Therefore, the guest OS in `guest.c` is a single stream of instructions, with no separation between the guest user space and the guest OS.



## assignment2

Updated automatically every 5 minutes



**Figure 1:** VM Run Loop in `simple-kvm.c`

The userspace hypervisor program `simple-kvm.c` uses the `/dev/kvm` file to create a virtual machine, allocate memory and CPU. The program then starts a KVM run loop that starts a vCPU using `KVM_RUN` and continuously handles the `VM_EXIT` routines when the vCPU returns from `KVM_RUN`, as shown in Figure 1. A simple print character operation by a guest translates to a write on port `0xE9` as a hypercall. The hypercall implementation for this setup uses the write IO operation to trigger a guest OS `VM_Exit`, and the exit cause is saved. The hypervisor checks that the VM exited trying to do `IO_OUT` operation on port `0xE9`, and then the userspace handler of the hypervisor prints the value written to the port.

The x86 architecture has different operating modes: real mode, protected mode, paged 32-bit, and long mode/paged 64-bit mode. Among other things, these modes refer to how the CPU will perform address translation. You can use the `simple-kvm.c` program to create and operate a VM in any of these



## assignment2

Updated automatically every 5 minutes

1. Go through the code and KVM APIs to understand the workings of the hypervisor. With the help of a well-drawn, labeled, and descriptive **flowchart**, explain the logical actions (**at least ten actions**) to setup and execute a VM in the **long mode of operation**.

Note that this step needs a flowchart diagram; no other output will be considered valid.

2. For each of the logical actions (**at least ten**) mentioned in the above flow chart, describe the KVM APIs associated with the actions and their purpose in your own words.

e.g.,

*1) Set up the physical memory for the guest OS... To assign the memory to the guest OS, the KVM\_YYY\_ZZZ ioctl call is used.*

*KVM\_YYY\_ZZZ:*

*Brief description of inputs to the ioctl call*

*Functionality that the ioctl call provides, etc.*

...

**Bonus:** If you want to learn more about a full hypervisor implementation, you can refer to kvmtool, a lighter hypervisor than QEMU. <https://github.com/kvmtool/kvmtool>

## 1b. Hypercalls: How can



## assignment2

Updated automatically every 5 minutes

the help of KVM. The user space handler can then determine the exit reason and direction of the data and handle it correspondingly.

**Note:** The boilerplate code already implements an example hypercall `HC_print8bit()`; which writes a byte value to the port `0xE9` using assembly instruction `outb`.

Implement the following hypercalls to extend your `simple-kvm.c` to be not-so-simple.

1. `void HC_print32bit(uint32_t val)`  
This hypercall should take a 32-bit value as input, and the hypervisor should handle the hypercall by printing the 32-bit value on the terminal. Make sure to use a single newline `"\n"` or `"endl"` after printing the 32-bit value inside the hypervisor.
2. `uint32_t HC_numExits()`  
The hypercall should return the number of times the guest VM has exited to the hypervisor from the beginning of its execution. The hypervisor should maintain the count and return it to the guest. Since the guest can't print directly on the terminal, please use `HC_print32bit()`, implemented in the previous step, to print the count.
3. `void HC_printStr(char *str)`



## assignment2

Updated automatically every 5 minutes

the hypervisor. To achieve this, you must pass the address of the string to the hypervisor, which will fetch the string from the guest's memory. Use `HC_numExits` to count the instances that exist before and after the hypercall. The count should change by one due to `HC_printStr()`.

**Note:** You should not add any extra new lines. If the string passed by the guest has a new line, it should be printed.

- The output of the simple-kvm should show the difference of two exits (one for `HC_numExits` and one for `HC_printStr`)

#### 4. `char *`

`HC_numExitsByType()`

The hypervisor should handle the hypercall by returning an address to a string of the following format:

*IO in: x*

*IO out: y*

The values `x` and `y` should be the actual values of VM exits due to IO in and IO out. The `char *` value should then be passed to `HC_printStr` for printing via hypercall.

#### 5. `uint32_t`

`HC_gvaToHva(uint32_t gva)`

The hypercall should return the Host Virtual Address (HVA) corresponding to a given Guest Virtual Address (GVA). If the guest asks to translate an invalid GVA for



## assignment2

Updated automatically every 5 minutes

HC\_print32bit to print the HVA value.

**Hint:** Lookup usage of the KVM\_TRANSLATE request flag and ioctl with the vmfd.

**Hint:**

- kvm\_run structure is heavily used in hypercalls.
- outl and in assembly instructions should be used for writing and reading 32-bit values from a serial port.

**Note:** To implement multiple hypercalls, you may use a different serial port for each hypercall or a struct with a field with a hypercall number and all arguments to the hypercall. In the second method, the struct will be known to the guest, and the hypervisor and the guest will pass the address to the struct for every hypercall with the correct hypercall number and the arguments set. The hypervisor can now read the hypercall number and the argument from the guest's memory. The implementation choice is left to you.

## 2. Build the matrix cloud

The architect (owner) of the matrix cloud generously offers users powerful bare-metal systems for their use. The architect was buying new machines for new users joining the matrix, and it became clear that many users were not using their powerful machines at maximum capacity. Inspired by insights gained in CS695, the architect decided to create a next-generation matrix cloud by moving users' systems to virtual machines, unlocking various advantages of virtualization.





## assignment2

Updated automatically every 5 minutes

the twist. Neo, under a special agreement (SLA) for high availability, claims 70% of the scheduling time (since he is the one), while Morpheus gets 30%.

The architect was able to create a starter code for the cloud but needs help to do the rest. Being a student of CS695, it's your task to help the architect achieve his dream.

Starter Code: `assignment2/part2/`

The starter code has `matrix.c`, which is the custom KVM hypervisor, sets up and runs two guest programs, `guest1.s` and `guest2.s`, inside the respective VMs. Each VM is set up with one vCPU each, and the startup code can configure them to run concurrently (via the Linux scheduler, scheduling each of the vCPU (pthread) threads independently).

Note that `matrix.c` is the control program and cannot proactively preempt or de-schedule the VMs. It has to wait for VM\_EXITs for control to reach back to it. The guest VMs have the following functionalities —

- `guest1.s`  
Issues a hypercall (using port 0x10) with the value in the `ax` register, and then on return from hypercall increment value stored in the `ax` register, and do this in an infinite loop.
- `guest2.s`  
Issues a hypercall (using port 0x11) with the value in the `ax` register, and then on return from hypercall decrement value stored in



## assignment2

Updated automatically every 5 minutes

that both VMs execute in **real** mode.

**Note: The userspace program uses *pthreads* to handle VM exits in separate user threads. These threads are not to be confused with per vCPU threads that Linux uses for scheduling the VMs.**

```
$ ./matrix
VMFD: 4,
Loaded
Program with
size: 7
VMFD: 5,
Loaded
Program with
size: 7
VMFD: 4
started
running
VMFD: 5
started
running
VMFD: 4
stopped
running -
exit reason:
2
VMFD: 4
KVM_EXIT_IO
VMFD: 4 out
port: 16,
data: 0
VMFD: 5
stopped
running -
exit reason:
2
VMFD: 5
KVM_EXIT_IO
VMFD: 5 out
port: 17,
data: 0
VMFD: 4
```

The output shows the output of two VMs running concurrently.

The output is slightly controlled with a sleep statement in the user program. Otherwise, the output would have been even more randomized.

[ The initial values of VMFD can be different other than 4,5, say x,y, but these x,y will be the same throughout a single run. But will stay consistent in runtime; this is applicable in all examples below]



Published using Google Docs

[Report abuse](#)[Learn more](#)

assignment2

Updated automatically every 5 minutes

```
running - exit
reason: 2
VMFD: 4
KVM_EXIT_IO
VMFD: 5
stopped
running - exit
reason: 2
VMFD: 5
KVM_EXIT_IO
VMFD: 5 out
port: 17,
data: 65535
VMFD: 4 out
port: 16,
data: 1
VMFD: 5
started
running
VMFD: 4
started
running
VMFD: 5
stopped
running - exit
reason: 2
VMFD: 5
KVM_EXIT_IO
VMFD: 5 out
port: 17,
data: 65534
VMFD: 4
stopped
running - exit
reason: 2
VMFD: 4
KVM_EXIT_IO
VMFD: 4 out
port: 16,
data: 2
.....
```

## 2a. One at a time

For your assignment, you must ensure the vCPU threads run on the



## assignment2

Updated automatically every 5 minutes

same using `htop`), modify the `kvm_run_vm` function to not make any `pthread` calls, but instead use the main thread to alternate between the two VMs on a `KVM_EXIT_IO`. This ensures that you can run two VMs on one physical CPU.

```
$ ./matrix
VMFD: 4,
Loaded
Program with
size: 7
VMFD: 5,
Loaded
Program with
size: 7
VMFD: 4
started
running
VMFD: 4
stopped
running -
exit reason:
2
VMFD: 4
KVM_EXIT_IO
VMFD: 4 out
port: 16,
data: 0
VMFD: 5
started
running
VMFD: 5
stopped
running -
exit reason:
2
VMFD: 5
KVM_EXIT_IO
VMFD: 5 out
port: 17,
data: 0
VMFD: 4
started
running
VMFD: 4
```

The single-threaded execution model ensures that the VM runs one after another. The execution model depends on the VM to ensure it returns to the userspace program after some time.

Note that in this sequence, the VMs strictly alternate and can be seen by switching the VMFD values in the corresponding outputs.



Published using Google Docs

[Report abuse](#)[Learn more](#)

assignment2

Updated automatically every 5 minutes

```
port: 16,
data: 1
VMFD: 5
started
running
VMFD: 5
stopped
running -
exit reason:
2
VMFD: 5
KVM_EXIT_IO
VMFD: 5 out
port: 17,
data: 65535
VMFD: 4
started
running
VMFD: 4
stopped
running -
exit reason:
2
VMFD: 4
KVM_EXIT_IO
VMFD: 4 out
port: 16,
data: 2
VMFD: 5
started
running
VMFD: 5
stopped
running -
exit reason:
2
VMFD: 5
KVM_EXIT_IO
VMFD: 5 out
port: 17,
data: 65534
.....
```

**Update:** If you have implemented part 2a in `matrix.c`, please make a copy of the



assignment2

Updated automatically every 5 minutes

a.s) to fit the implementation.

**Please make sure that `matrix-a.c` should only implement part 2a!**

**This checkpoint should ensure that you have completed 2a**

**2b. Vulnerabilities found... Please fix it.**

Remove/Comment out the `out %ax, $0x10` and `out %ax, $0x11` in `guest1.s` and `guest2.s`. You can try to run your VMs now but will observe that only `vm1` (or whatever VMs you were running first) will be running. Since no IO operation is happening, no VM exits will happen. Therefore, the previous mechanism will not work. A sample output is as follows.

```
$ ./matrix
VMFD: 4,
Loaded
Program with
size: 7
VMFD: 5,
Loaded
Program with
size: 7
VMFD: 4
started
running
```

The other VM will starve if the running VM doesn't return to the user space program, which is happening in the example.

With this setup, how do we achieve the same control as in **2a** when no I/O is happening? 🤔

The Linux scheduler uses timer interrupts to schedule processes; therefore, a similar approach can also be used to schedule the VMs. The `timer_create()` subroutine of the time library creates a



## assignment2

Updated automatically every 5 minutes

for the signal to be injected periodically.  
Once this is set – the timer ready to fire — attempt to run your VMs. Only one of them should run. On timer expiry, a signal is delivered to the process, which will run the guest VM. On such a signal, the default behavior should result in a VM exit to be handed by the user space program.

This also requires the proper setup of signals within the KVM. You must use the KVM  
API `KVM_SET_SIGNAL_MASK` to do the same during VM execution. Set the mask to clear the signal that you want to be delivered. Additionally, make sure that the signal is blocked for the control thread.

For sanity check, add the following line after `ioctl` return  
from `KVM_RUN`:  
`printf("Time: %f\n",  
CURRENT_TIME);`

If everything is done correctly, the output will be something like that listed below —

```
$ ./matrix
VMFD: 4,
Loaded
Program with
size: 5
VMFD: 5,
Loaded
Program with
size: 5
VMFD: 4
started
running
Time:
1.002131
VMFD: 4
stopped
```

If signals are injected into the VMs, the execution will be handled by the user space program, which can decide who to run next. In this example, the userspace runs one VM after another. For the example, the `QUANTUM` was set at 1s.



Published using Google Docs

[Report abuse](#)

[Learn more](#)

assignment2

Updated automatically every 5 minutes

running Time: 2.002100 VMFD: 5 stopped running - exit reason: 10 VMFD: 5 KVM_EXIT_INTR VMFD: 4 started running Time: 3.002045 VMFD: 4 stopped running - exit reason: 10 VMFD: 4 KVM_EXIT_INTR VMFD: 5 started running Time: 4.002001 VMFD: 5 stopped running - exit reason: 10 VMFD: 5 KVM_EXIT_INTR VMFD: 4 started running Time: 5.001990 VMFD: 4 stopped running - exit reason: 10 VMFD: 4 KVM_EXIT_INTR
--





Published using Google Docs

[Report abuse](#)

[Learn more](#)

assignment2

Updated automatically every 5 minutes

```
running -
exit reason:
10
VMFD: 5
KVM_EXIT_INTR
VMFD: 4
started
running
Time:
7.001829
VMFD: 4
stopped
running -
exit reason:
10
VMFD: 4
KVM_EXIT_INTR
VMFD: 5
started
running
Time:
8.001799
VMFD: 5
stopped
running -
exit reason:
10
VMFD: 5
KVM_EXIT_INTR
VMFD: 4
started
running
Time:
9.001752
VMFD: 4
stopped
running -
exit reason:
10
VMFD: 4
KVM_EXIT_INTR
VMFD: 5
started
running
.....
```



## assignment2

Updated automatically every 5 minutes

needs careful setup.

**Update:** If you have implemented part 2b in `matrix.c`, please make a copy of the implementation in a file named `matrix-b.c` and update the makefile to generate `matrix-b` executable accordingly. You should also make separate guests (`guest1-b.s` & `guest2-b.s`) to fit the implementation.

**Please make sure that `matrix-b.c` should only implement part 2b!**

**This checkpoint should ensure that you have completed 2b**

## 2c. The final leap

Since you can now control the order of execution of the VMs, implement fractional scheduling for extending solution of 2b, which schedules `vm1` seven times and `vm2` three times out of ten time-quantum. Define the time quantum as a macro (`QUANTUM`) with a value of **1 second** (no matter what unit you are using, the value should equal 1s) and assume there will be no I/O operations in your VMs. Create two more macros, `FRAC_A` and `FRAC_B`, for scheduling fractions of `vm1` and `vm2` out of ten. By default, it should be seven and three.

References:

- [https://man7.org/linux/man-pages/man2/timer\\_create.2.html](https://man7.org/linux/man-pages/man2/timer_create.2.html)
- <https://docs.kernel.org/virt/kvm/api.html>

**Update:** No need to make a separate file for this section.



assignment2

Updated automatically every 5 minutes

[work with matrix.c.](#)

## Submission instructions

- Your code should be well-commented and readable.
- You must submit a PDF file for Part 1a [named part1a.pdf] and final code solutions for Parts 1b and 2 according to the directory structure. Make sure that you perform a `make clean` before submitting solutions.
- All submissions have to be done via Moodle.  
Name your submission folder as  
`<rollnumber>_assignment2` and the tarball as  
`<rollnumber>_assignment2.tar.gz`  
(e.g.  
`22d0999_assignment2.tar.gz`)  
- all small letters

**Please strictly adhere to this format; otherwise, your submission will not count.**

`22d0999_assignment2.gz`  
`22d0999_assignment2.zip`  
`22d0999_assignment2.tar.xz`

- *You can create the tarball using (copy-paste this command for a happy life)*  
`tar -czvf`  
`<rollnumber>_assignment2.tar.gz`  
`<rollnumber>_assignment2`
- Before submission, make



## assignment2

Updated automatically every 5 minutes

- Modify the README.md files to fit your submission.

- The tar should contain the following files in the following directory structure:

```
<rollnumber>_assignment2/
```

```
├── .git/
│   └── ..
. /* all git-related
files */
├── part1b/
│   ├── guest.c
│   ├── guest.ld
│   ├── guest16.s
│   ├── Makefile
│   ├── payload.ld
│   ├── README.md
│   └── simple-kvm.c
├── part2/
│   ├── guest1.s
│   ├── guest1-a.s
│   ├── guest1-b.s
│   ├── guest2.s
│   ├── guest2-a.s
│   ├── guest2-b.s
│   ├── Makefile
│   └── matrix-a.c
```



Published using Google Docs

[Report abuse](#)

[Learn more](#)

assignment2

Updated automatically every 5 minutes

```
├── .gitignore
├── Makefile
└── README.md
```

**Deadline: 19th**  
**Feb 2024, 11:59**  
**PM via Moodle.**  
**21st**  
**Feb 2024**  
**(extended)**